

502: Operating Systems - Assessed Coursework

Producer-Consumer

Due date: **8th December 2017**
(Submission on CATe by 18:59)

The purpose of this exercise is to use the producer-consumer scenario to gain experience into using system calls for semaphores. A typical producer-consumer scenario consists of a group of producers 'producing' jobs with a group of consumers 'consuming' the jobs. In order for the producers and consumers to be able to communicate with each other, a shared data structure is traditionally used to store the jobs and with which the producers and consumers interact. Due to the shared nature of this data structure, it needs to be protected by synchronisation mechanisms to ensure consistent access to the data.

In this exercise, we implement the producer-consumer scenario using POSIX threads, where each producer and consumer will be running in a different thread. The shared data structure used will be a circular queue, which can be implemented as an array for the purpose of this exercise. This array will contain details of the job, where a job has a job id (which is one plus the location that they occupy in the circular queue) and a duration (which is the time taken to 'process' the job - for the purpose of this exercise, processing a job will mean that the consumer thread will sleep for that duration).

We will implement a **main** program that will contain two functions - one for the producer and one for the consumer, each of which will be run in a separate thread. Depending on the number of producers and consumers required, we will be creating those many threads.

The sequence of steps for the main program (and the two functions) are as follows:

1. Main programme
 - (a) Read in four command line arguments - size of the queue, number of jobs to generate for each producer (each producer will generate the same number of jobs), number of producers, and number of consumers.
 - (b) Set-up and initialise required variables, as necessary.
 - (c) Set-up and initialise semaphores, as necessary.
 - (d) Create the required producers and consumers.
 - (e) Quit, but ensure that there is process clean-up.
2. Producer
 - (a) Initialise parameters, as required.
 - (b) Add the required number of jobs to the queue, with each job being added once every 1 to 5 seconds. If a job is taken (and deleted) by the consumer, then another job can be produced which has the same id. Duration for each job should be in the range 1 to 10 seconds. If the circular queue is full, block while waiting for an empty slot and if a slot doesn't become available after *20 seconds*, quit, even though you have not produced all the jobs.
 - (c) Print the status (example format given in *example_output.txt*).
 - (d) Quit when there are no more jobs left to produce.
3. Consumer
 - (a) Initialise parameters, as required.
 - (b) Take a job from the circular queue and 'sleep' for the duration specified. If the circular queue is empty, block while waiting for jobs and quit if no jobs arrive within *20 seconds*.

- (c) Print the status (example format given in *example_output.txt*).
- (d) If there are no jobs left to consume, wait for *20 seconds* to check if any new jobs are added, and if not, quit.

Your solution must handle multiple producers and consumers concurrently and use semaphores to restrict access to the shared resource (in this case, the circular queue). The solution should also deal with input errors, system call errors, and any other errors, appropriately. You are welcome to come up with challenging/complicated/interesting solutions ☺. *Hint:* In order to make the Producer/Consumer wait for 20 seconds, the elegant (and preferred) solution is to implement this delay as part of the semaphore. Useful command-line tools for debugging by checking the status of semaphores are *ipcs* and *ipcrm*.

Notes: Download the outline source code from CATe. You are not obliged to use these; however they should simplify the process of achieving working solutions. You can choose to change these files as necessary. A few notes on the layout and support files follow:

- These support files have been tested on Linux in the DoC labs. You can however choose to use your own programming environment.
- The file *main.cc* contains the overall structure of the programme and has a simple example of creating and using a thread.
- The file *helper.h* includes the required header files, as well as some of the pre-defined variables. Please ensure that you change the value of the key used for the semaphore in-order for it to be unique amongst students.
- The file *helper.cc* contains a few helper functions, including functions for creating, using and destroying semaphores.
- The file *example_output.txt* contains an example output using the following parameters - Queue size = 5, Number of Items produced per producer = 6, Number of producers = 2, and Number of consumers = 3. Please note this is just a suggested output - you do not have to make your output look *exactly* similar ☺.
- The *Makefile* allows users to use the *make* command on Linux to compile the various parts of the exercise. It can also be used to help configure your preferred development environment with the correct commands, flags and parameters.

You should hand in the following on CATe:

1. An archive (**code.tgz**) containing **only** the C++ files (*main.cc*, *helper.cc*, *helper.h*) and the *Makefile*. Please remove any extraneous directories (such as *.git*, *__MACOSX*, etc.)

Tutorials

Please be aware that the basic functionalities of these may slightly differ from the lectures.

Semaphores: <http://beej.us/guide/bgipc/output/html/multipage/semaphores.html>

POSIX Threads Programming: <https://computing.llnl.gov/tutorials/pthreads/>

Useful Links

Workstations in DoC: <https://www.doc.ic.ac.uk/csg/facilities/lab/workstations>

Remote Login in DoC: <https://www.doc.ic.ac.uk/csg/services/linux>

Creating and Unpacking .tar.gz files: <http://arxiv.org/help/tar>

ipcs - provide information on ipc facilities: <http://linux.die.net/man/1/ipcs>

ipcrm - remove a message queue, semaphore set or shared memory id:

<http://linux.die.net/man/1/ipcrm>